



# Instructions for Crème Brûlée

Matthias Funke  
Professional Services

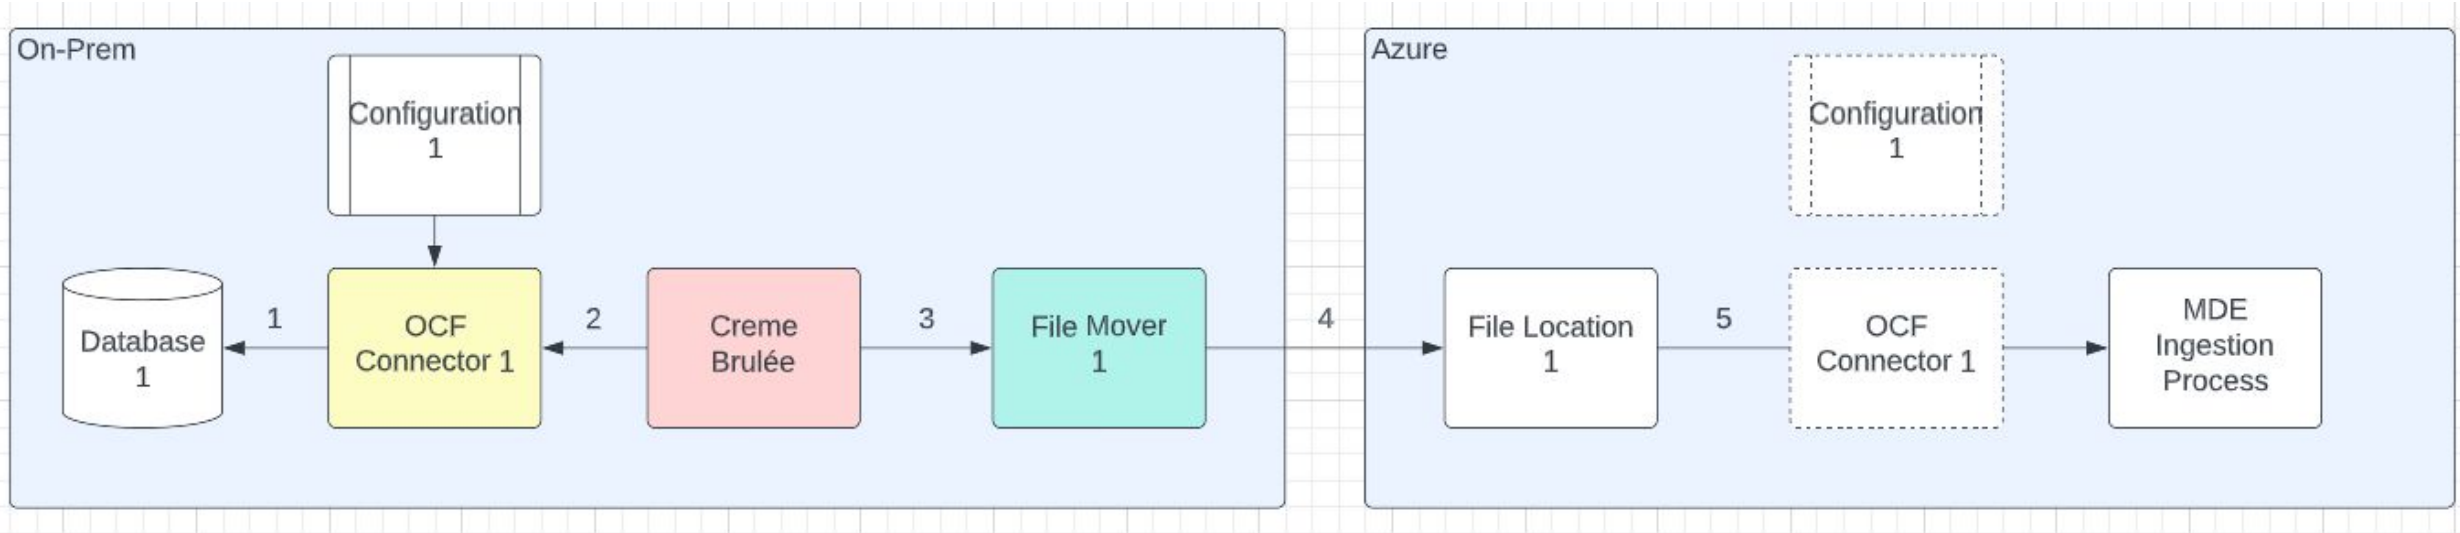
# Allow MDE remotely

- Vattenfall wants to be able to do MDE remotely
- Solution described here uses existing OCF connectors (should work for any RDBMS OCF connector)
- Tested with SQL Server, Snowflake, Athena, Glue, Postgres
- Custom code "creme\_brulee.py" is lightweight and low maintenance:
  - filters out blacklisted schemas
  - converts file format to make suitable for ingestion

# Vattenfall Remaining Work

- Build framework for moving files
- Build framework for deploying containers automatically for each datasource
- Build framework for config management:
  - JDBC URI
  - username
  - password
  - any other params required for some

# End-to-End Process



1. Yellow container uses OCF framework (built in Java)
2. Red container uses Python to transform the files into suitable format and filters out undesired schemas
3. File Mover (Vattenfall) moves the files to the Alation server
4. Ingestion on Alation server is done using shadow OCF connector (file based)

# OCF Container

1. Obtain the connector ZIP file from Alation
2. Unzip it
3. Load the image
4. Verify the image

```
sh-4.2$ aws s3 cp s3://com.alationpro/creme/alation-awsglue-1.0.1.1117.zip .
```

```
download: s3://com.alationpro/creme/alation-awsglue-1.0.1.1117.zip to
```

```
./alation-awsglue-1.0.1.1117.zip
```

```
sh-4.2$ unzip ./alation-awsglue-1.0.1.1117.zip
```

```
Archive:  ./alation-awsglue-1.0.1.1117.zip
```

```
  inflating: alation-awsglue-ocf:1.0.1.1117.img
```

```
  inflating: MANIFEST.MF
```

```
sh-4.2$ sudo docker image load -i alation-awsglue-ocf:1.0.1.1117.img
```

```
Loaded image: alation-awsglue-ocf:1.0.1.1117
```

```
sh-4.2$ sudo docker image ls |grep glu
```

```
alation-awsglue-ocf          1.0.1.1117          89119871006a    4 months ago    252MB
```

# Create a config.json file

## 1. The relevant keys are listed in the MANIFEST file

```
[
  {
    "key": "jdbc_uri",
    "value": "sqlserver://sql-server.alationproserv.com:1433",
    "type": "TEXT"
  },
  {
    "key": "use_ntlm",
    "value": "True",
    "type": "BOOLEAN"
  },
  {
    "key": "domain_name",
    "value": "ALATTIONPROSERV",
    "type": "TEXT"
  },
  {
    "key": "username",
    "value": "alation-svc",
    "type": "TEXT"
  },
  {
    "key": "password",
    "value": "<REDACTED>",
    "type": "ENCRYPTEDTEXT"
  }
]
```

# Create a my\_conf.yaml file

1. Blacklist any schemas you don't want (all lower case)
2. Provide data source ID (optional)
3. Provide file name for CSV file (optional)

```
# yaml file for snowflake MDE
blacklist:
- citibike.mytestschema
- citibike.mytestschema2

ds_id: 20
csv_file_name: my_csv.csv
```

# Create shared Docker volume

1. RAM disk means no passwords are permanently saved to disk

```
docker volume create --driver local --opt type=tmpfs --opt device=tmpfs --opt o=size=100m,uid=1000 foo
```

```
[root@ip-172-31-5-116 _data]# docker volume inspect foo
[
  {
    "CreatedAt": "2022-12-12T13:39:47Z",
    "Driver": "local",
    "Labels": {},
    "Mountpoint": "/var/lib/docker/volumes/foo/_data",
    "Name": "foo",
    "Options": {
      "device": "tmpfs",
      "o": "size=100m,uid=1000",
      "type": "tmpfs"
    },
    "Scope": "local"
  }
]
```



# Copy files into Docker volume

```
[root@ip-172-31-5-116 _data]# ls -l
total 28
-rw-r--r-- 1 root root 4577 Dec 12 10:18 creme_brulee.py # code for filtering
-rw-r--r-- 1 root root  156 Dec 12 10:17 my_conf.yaml # configuration for filters
-rw-r--r-- 1 root root  326 Dec  6 11:27 snow.json # configuration for data source
```

# Start OCF Container

1. Start the container as a daemon
2. Verify it has started

```
[root@ip-172-31-5-116 _data]# docker run -d -e PORT=8980 -v foo:/creme --name alation-awsglue-ocf
alation-awsglue-ocf:1.0.1.1117
825d31e72941872ea76a8e1dfc75bc5d803a455134f18128eebe3a1f2a3eb8da
[root@ip-172-31-5-116 _data]# docker logs $(docker ps -f name=alation-awsglue-ocf -q)
INFO    2022-12-12 13:51:06,351 [main] alation.sdk.core.grpc.server.GrpcServer - Attempting to start
the Server
INFO    2022-12-12 13:51:06,462 [main] alation.sdk.core.grpc.server.GrpcServer - Server started,
listening on 8980
[root@ip-172-31-5-116 _data]#
```

# Run MDE in OCF Container

The command starts MDE in the correct container with the correct configuration

```
[root@ip-172-31-5-116 _data]# docker exec $(docker ps -f name=alation-awsglue-ocf -q) java -cp /opt/APP.jar alation.sdk.rdbms.grpc.client.AbstractRDBMSClient -p 8980 --conf /creme/glue.json -o /creme CONF
INFO      2022-12-12 14:09:02,841 [main] alation.sdk.rdbms.grpc.client.ConfigurationVerificationClient - Send request to verify configuration!
INFO      2022-12-12 14:09:06,436 [main] alation.sdk.core.grpc.client.AbstractClient - success

[root@ip-172-31-5-116 _data]# docker exec $(docker ps -f name=alation-awsglue-ocf -q) java -cp /opt/APP.jar alation.sdk.rdbms.grpc.client.AbstractRDBMSClient -p 8980 --conf /creme/glue.json -o /creme META
INFO      2022-12-12 14:09:12,834 [main] alation.sdk.rdbms.grpc.client.MetadataExtractionClient - Send request to run metadata extraction!
INFO      2022-12-12 14:09:18,148 [main] alation.sdk.core.grpc.client.AbstractClient - success
[root@ip-172-31-5-116 _data]#
```

# Optional: inspect output files

Successful MDE produces .txt files

```
[root@ip-172-31-5-116 _data]# ls -lrt
```

```
total 1000
```

```
-rw-r--r-- 1 root root    326 Dec  6 11:27 snow.json
-rw-r--r-- 1 root root    296 Dec  7 12:42 analytics.json
-rw-r--r-- 1 root root    156 Dec 12 10:17 my_conf.yaml
-rw-r--r-- 1 root root   4577 Dec 12 10:18 creme_brulee.py
-rw-r--r-- 1 root root    292 Dec 12 14:08 glue.json
-rw-r--r-- 1 root root  33718 Dec 12 14:09 table.txt
-rw-r--r-- 1 root root   4002 Dec 12 14:09 schema.txt
-rw-r--r-- 1 root root   1097 Dec 12 14:09 extractionProgress.txt
-rw-r--r-- 1 root root 949877 Dec 12 14:09 column.txt
-rw-r--r-- 1 root root      8 Dec 12 14:09 Status.txt
```

```
[root@ip-172-31-5-116 _data]# cat schema.txt
```

```
{"id_":{"name_":{"original_":"default","normalized_":"","memoizedIsInitialized":-1,"unknownFields":{"fields":{},"fieldsDescending":{}},"memoizedSize":-1,"memoizedHashCode":0},"memoizedIsInitialized":-1,"unknownFields":{"fields":{},"fieldsDescending":{}},"memoizedSize":-1,"memoizedHashCode":0},"sourceComment_":"Default Hive database","memoizedIsInitialized":-1,"unknownFields":{"fields":{},"fieldsDescending":{}},"memoizedSize":-1,"memoizedHashCode":0}
```

# Build Creme Brulee container from Dockerfile

Place Dockerfile and requirements.txt in a separate directory

```
[root@ip-172-31-5-116 creme_brulee]# docker build . -t creme-brulee:1.2
```

```
Sending build context to Docker daemon 8.192kB
```

```
Step 1/4 : FROM python:3.8
```

```
..
```

```
Removing intermediate container b3743e802c67
```

```
----> a4eea27eed0a
```

```
Successfully built a4eea27eed0a
```

```
Successfully tagged creme-brulee:1.2
```

```
[root@ip-172-31-5-116 creme_brulee]# ls
```

```
Dockerfile requirements.txt
```

```
[root@ip-172-31-5-116 creme_brulee]# docker image ls
```

REPOSITORY	TAG	IMAGE ID	CREATED	SIZE
creme-brulee	1.2	a4eea27eed0a	2 minutes ago	1.08GB

# Convert txt files to dump files and filter schemas

Successful conversion produces .dump files. Note file names change

```
[root@ip-172-31-5-116 _data]# docker run -v foo:/creme creme-brulee:1.2 python3  
/creme/creme_brulee.py /creme/my_conf.yaml
```

All done.

```
[root@ip-172-31-5-116 _data]# ls -lrt
```

total 1332

```
-rw-r--r-- 1 root root    326 Dec  6 11:27 snow.json  
-rw-r--r-- 1 root root    296 Dec  7 12:42 analytics.json  
-rw-r--r-- 1 root root    156 Dec 12 10:17 my_conf.yaml  
-rw-r--r-- 1 root root   4577 Dec 12 10:18 creme_brulee.py  
-rw-r--r-- 1 root root    292 Dec 12 14:08 glue.json  
-rw-r--r-- 1 root root   33718 Dec 12 14:09 table.txt  
-rw-r--r-- 1 root root   4002 Dec 12 14:09 schema.txt  
-rw-r--r-- 1 root root   1097 Dec 12 14:09 extractionProgress.txt  
-rw-r--r-- 1 root root 949877 Dec 12 14:09 column.txt  
-rw-r--r-- 1 root root      8 Dec 12 14:09 Status.txt  
-rw-r--r-- 1 root root    672 Dec 12 14:14 Schema.dump  
-rw-r--r-- 1 root root   9849 Dec 12 14:14 Table.dump  
-rw-r--r-- 1 root root 281954 Dec 12 14:14 Attribute.dump  
-rw-r--r-- 1 root root  38448 Dec 12 14:14 my_csv.csv
```

/creme is the mounted volume  
inside the creme brulee container

# Move files to Alation Server (Vattenfall)

```
ls -l /data/site_data/tmp/rosemeta/
total 28
drwxrwxr-x 4 alation alation 4096 Dec  5 14:36 18
drwxrwxr-x 4 alation alation 4096 Dec  6 13:48 20
drwxrwxr-x 3 alation alation 4096 Dec  6 13:04 21
drwxr-xr-x 3 root      root      4096 Dec 12 14:30 28
drwxr-xr-x 3 alation alation 4096 Nov 29 12:49  4
drwxr-xr-x 5 alation alation 4096 Dec  6 12:51  6
drwxr-xr-x 2 alation alation 4096 Oct 24 09:17  7
```

```
sh-4.2$ find .
```

```
.
./6
./6/1969
./6/1969/Schema.dump
./6/1969/Catalog.dump
./6/1969/Table.dump
./6/1969/Attribute.dump
./6/2013
./6/2013/Schema.dump
./6/2013/Table.dump
```

These are data source IDs  
Use the correct OCF connector  
No other config required or possible

Directory name can be any number  
MDE will use the highest number  
Recommended: 4-digit

Example: 2013 is used, not 1969

# Alation Side: "Fake" OCF Connector

om/data/18/settings/extraction/

All Domains ▾ Search Alation

Sources Manage Settings / fake SQL server

To finish configuring your database, you will need to run a couple more actions.

1. Configure The Connection

You may need to configure the connection to your database before running other jobs.

To configure your connection, click on the General Settings tab, fill out any configuration options that are required for your database, and click "Test" under Test Connection to verify Alation can reach your database.

2. Run Data Sampling

Sampling your database generates content previews for tables found in your database.

To run Data Sampling, click on the Data Sampling tab, then click the 'Sample' button.

3. Run Query Log Ingestion

Query Log Ingestion imports your query history into Alation.

To preview Query Log Ingestion, click on the Query Log Ingestion tab, then click the 'Preview' button.

To run Query Log Ingestion, click on the Query Log Ingestion tab, then click the 'Import' button.

Access

General Settings

Metadata Extraction

Compose

Data Sampling

Per-Object Parameters

Query Log Ingestion

Custom Settings

Automated Extraction

A background job periodically fetches the database's metadata (including schemas, tables, columns, and procedures/functions) to keep our representation of the database up to date. You can schedule this job according to how frequently the database changes its structure.

The average run-time of this job is 0.45 seconds.

Application Settings

Enable Raw Dump Or Replay ⓘ


Enable Ingestion Replay ▾

Save

Ingestion Replay picks up the files from the previous step

Cannot use selective extraction here

Data Intelligence +





# Success Page

## Metadata Syncing Info

Metadata Syncing  
Started Dec 5 at 2:37pm, took 1 seconds.

### Status: Succeeded

RDBMS extraction/ingestion completed in 0.43 seconds, extracted all schemas.

### Details

rosemeta.DatabaseCatalog:

3 new object(s) created

rosemeta.Schema:

4 existing object(s) updated

rosemeta.Table:

12 existing object(s) updated

rosemeta.Attribute:

86 existing object(s) updated

FkFixer Info Log:

Fixing foreign keys for Schema, Tables, and Attributes finished in 0.02s.

StatsAggregator Info Log:

Counting tables per schema and attributes per table finished in 0.01s.

MetadataChangelog Info Log:

Creating metadata changelog finished in 0.01s.

SearchNotifier Info Log:

Notifying Search finished in 0.05s...

[Show More](#)

## Job Errors (0 in total)

*No data to display*

You can schedule MDE to run every  
hour / every day

No need for config changes